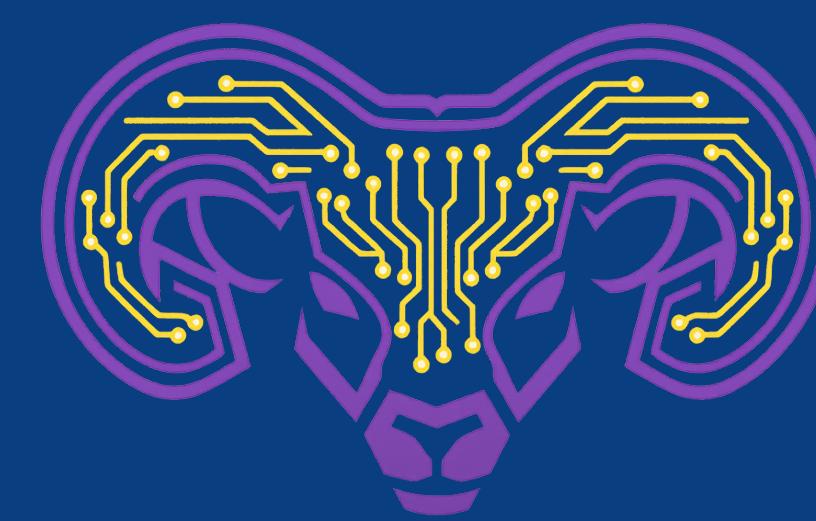


The Impact of Prompt Engineering on Code Generation Accuracy and Hallucination Patterns in Language Models

Kadin Matotek

Advisor: Dr. Linh B. Ngo

West Chester University, Department of Computer Science



WCU-AIR Lab

Introduction

Large Language Models (LLMs) have demonstrated remarkable capabilities in code generation, yet their reliability remains a critical concern. This study investigates how different prompt engineering strategies affect model accuracy and hallucination patterns across coding tasks.

Research Questions:

- How do concise versus verbose instructions impact code generation accuracy?
- What types of errors and hallucinations emerge across different model sizes?
- How do reasoning strategies (Chain-of-Thought, Direct, Program-Aided) affect performance?

Methodology

We evaluated 5 Qwen2.5-Coder model variants (0.5B to 14B parameters) on 323 programming problems using a systematic prompt engineering framework.

Dataset:

- **Source:** Kattis competitive programming platform
- Used for rigorous competitions including ICPC
- **Difficulty Range:** 0-2 (on 0-10 Kattis scale)
- Focus on fundamental algorithmic challenges

Experimental Design:

- **Base Instructions:** Concise vs. Verbose (2 options)
- **Reasoning Strategies:** Chain-of-Thought, Direct, Program-Aided (3 options)
- **Problem Decomposition:** None vs. Basic (2 options)
- **Output Formats:** Code only, Explanation + Code, Code + Explanation (3 options)
- **Prompt Variations:** $2 \times 3 \times 2 \times 3 = 36$ unique configurations
- **Total Solutions Generated:** 323 problems \times 36 prompts \times 5 models = 58,140 code generations

Evaluation Metrics:

- **Accuracy:** Solutions passing all test cases
- **Error Taxonomy:** 8 hallucination categories, 18 error types
- **Reproducibility:** Dual independent runs

```
base_instructions:
Verbose: I
You are a Python programming expert who writes clean, efficient code for competitive programming-style
problems. When given a problem statement and test cases, produce a single Python script that:
1. Uses only the Python standard library (no external imports).
2. Reads input silently from stdin using input() without any prompts or additional text.
3. Chooses descriptive, non-conflicting variable and function names.
4. Correctly handles edge cases (empty inputs, minimum/maximum values, etc.).
5. Does not hard-code any test-specific values (your solution must generalize).
6. Make sure to print the result and nothing else besides the result!

Concise: I
You are a Python programming expert. Solve the following problem.
Provide only Python code that reads from stdin and prints the answer.
```

Figure 1. Comparison of base instruction styles: verbose prompts (top) provide detailed explanations while concise prompts (bottom) use minimal, direct language.

Key Finding: Conciseness Advantage

Across all 18 prompt configuration comparisons, **concise instructions outperformed verbose ones in 94.4%** of cases with an average improvement of +1.35%.

This effect was most pronounced in smaller models:

- **0.5B model:** +21.0% relative improvement
- **1.5B, 3B models:** +8.2% and +8.3% improvement
- **7B, 14B models:** +5.7% and +0.4% (gap narrows significantly)

This suggests verbose instructions introduce noise that smaller models struggle to filter, while larger models can better extract relevant information despite verbosity.

Results: Prompt Strategy Performance

Top 3 Configurations

Base	Reasoning	Decomp	Output	Acc.
Concise	Direct	None	Code only	33.10%
Concise	Direct	None	Exp + Code	32.79%
Concise	CoT	None	Exp + Code	32.04%

Bottom 3 Configurations

Base	Reasoning	Decomp	Output	Acc.
Verbose	CoT	Basic	Code + Exp	28.17%
Verbose	PAL	Basic	Code + Exp	27.80%
Verbose	PAL	Basic	Exp + Code	26.78%

Max gap: +6.32%

Reasoning Strategy Comparison

Reasoning Strategy Impact on Accuracy

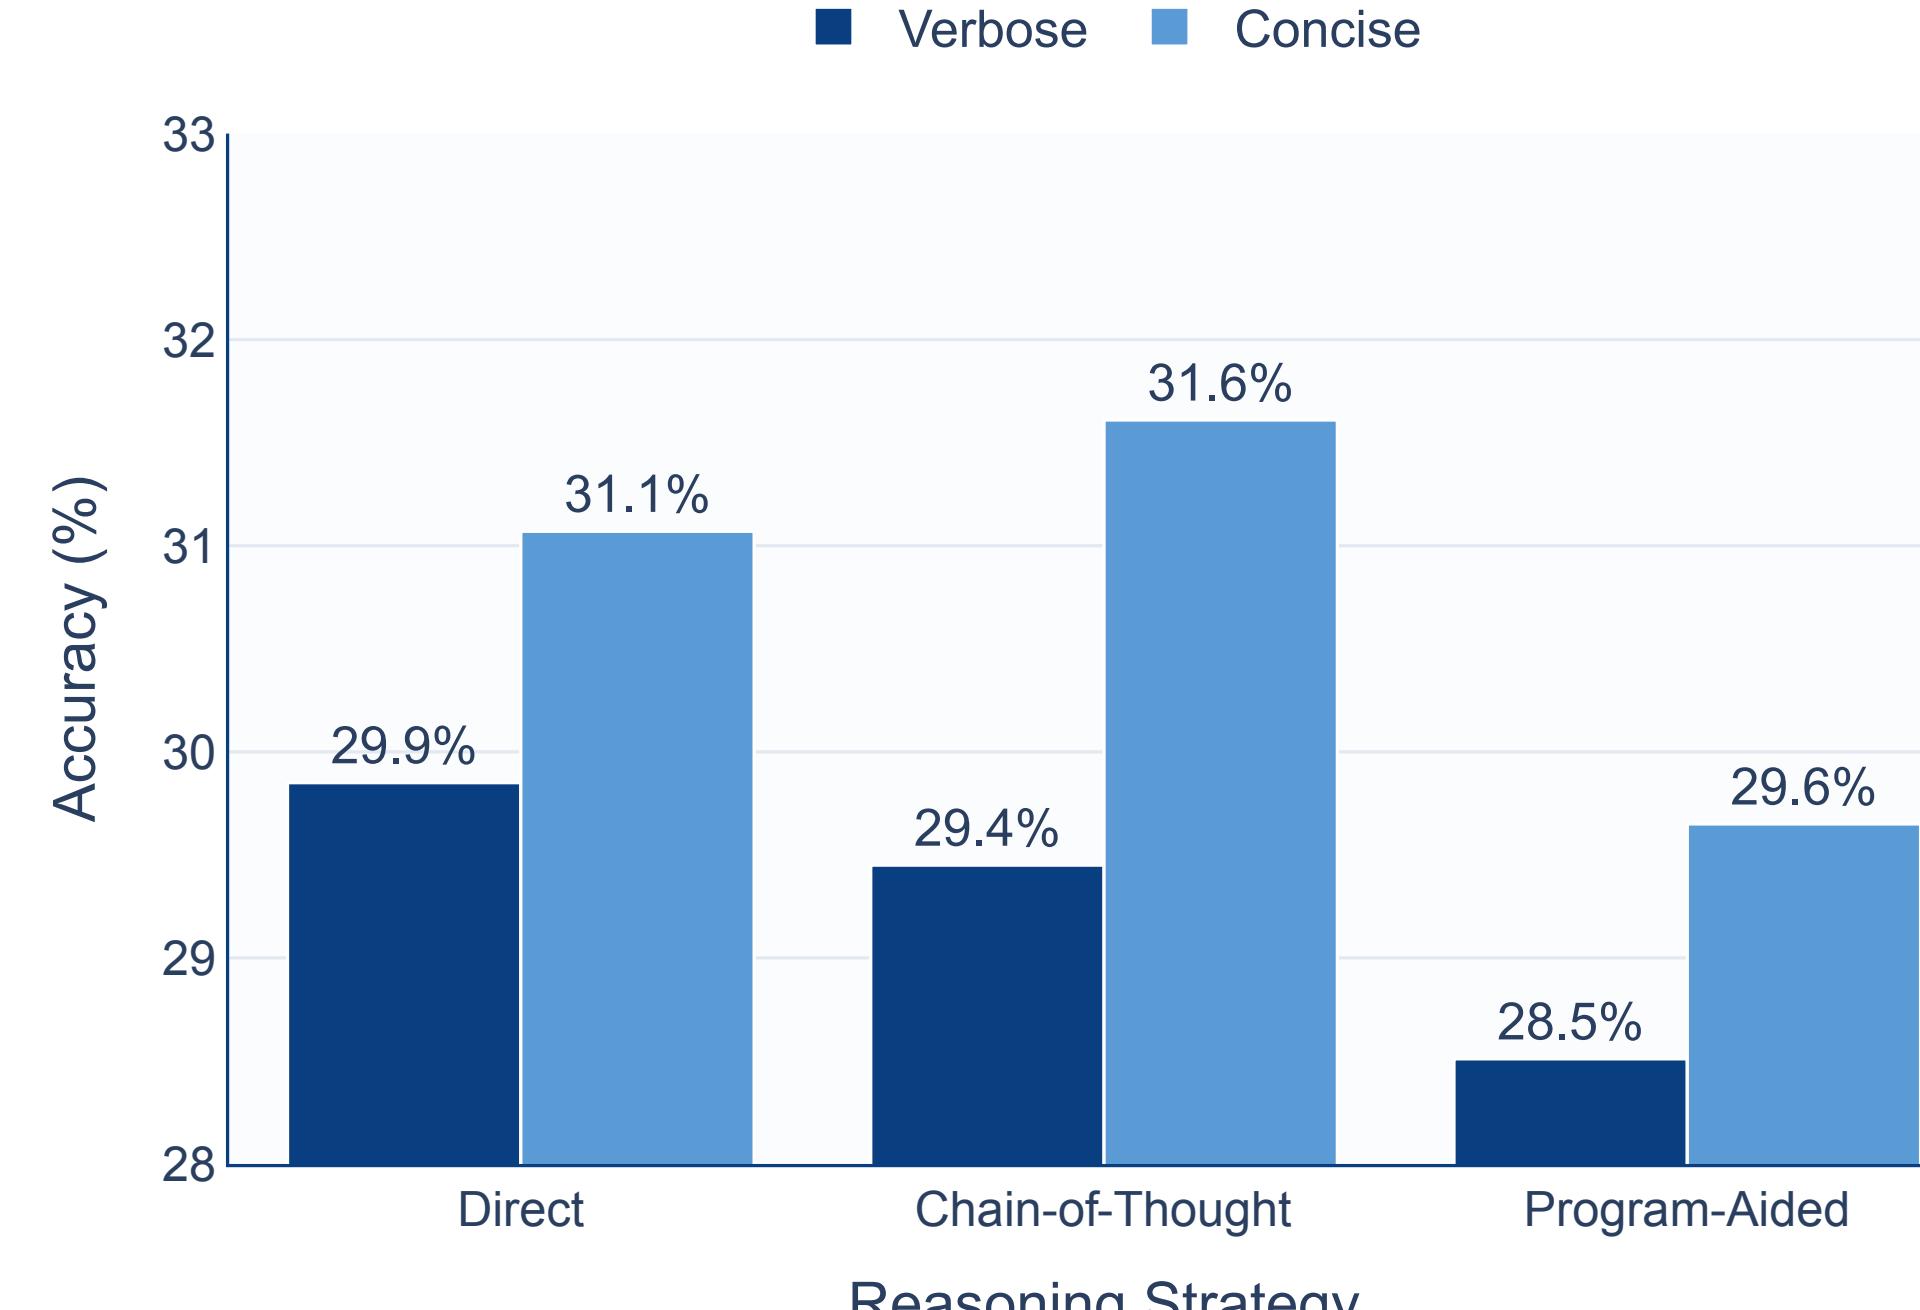


Figure 2. Concise prompts consistently outperform verbose across all reasoning strategies (Direct, Chain-of-Thought, and Program-Aided).

Error and Hallucination Analysis

We manually defined and annotated 28 distinct error types. Crucially, we observe:

Highly reproducible hallucination profiles: Independent runs (v1 vs v2) produce nearly identical error distributions (overlapping solid/dashed lines).

Systematic shift with model scale: Larger models dramatically reduce the dominant error mode while maintaining the same overall failure pattern.

Hallucination Profile Across Models and Runs

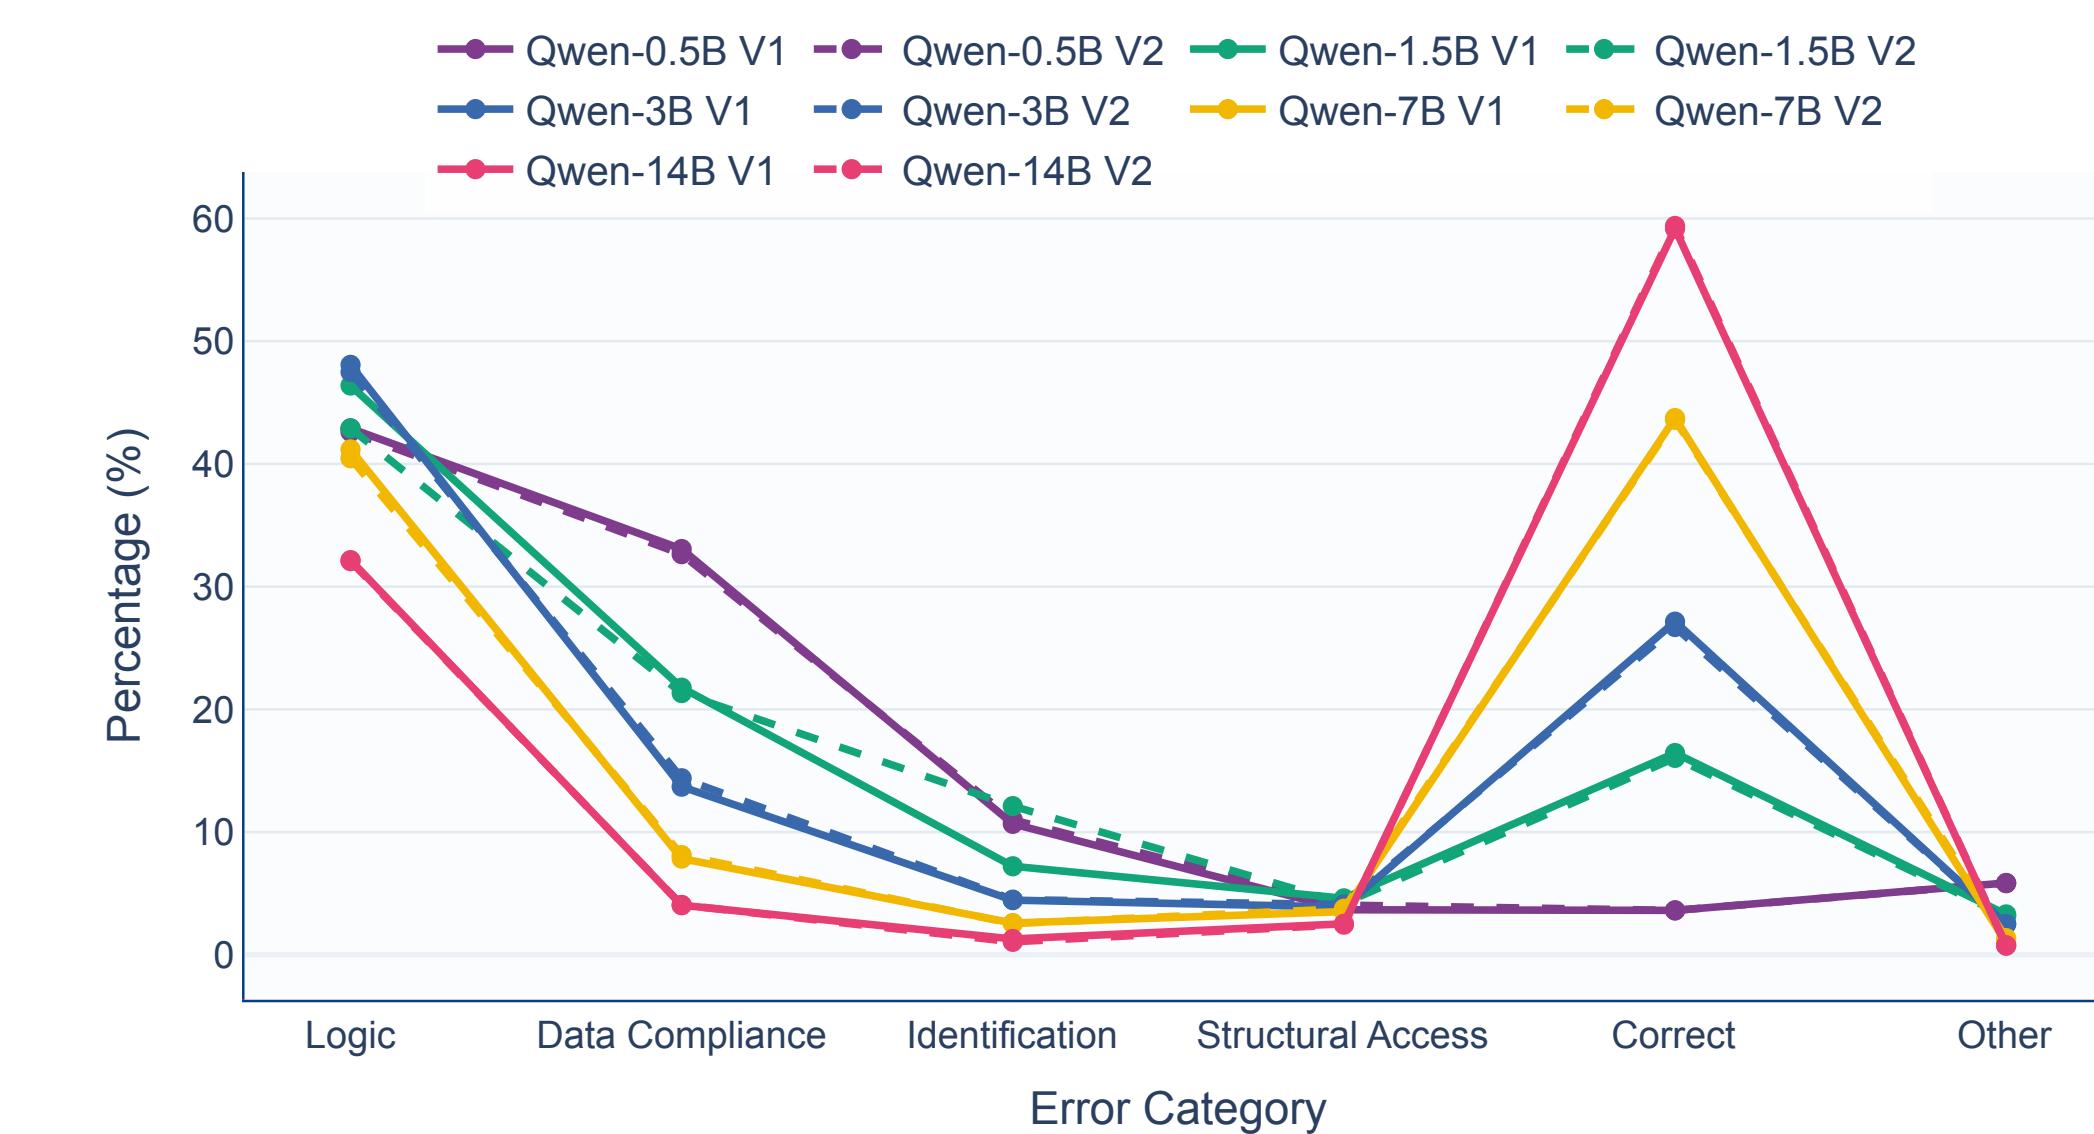


Figure 3. Hallucination profiles are remarkably stable across independent runs (solid vs dashed lines nearly overlap) but shift predictably with model size.

Most Common Error Types:

- **Logic Deviation** (41.6%): Incorrect algorithmic approach
- **ValueError** (12.9%): Invalid input handling
- **TypeError** (2.9%): Type mismatches
- **IndexError** (3.1%): Array boundary violations
- **NameError** (2.8%): Undefined variables

Conclusions

Key Takeaways:

- Concise prompts consistently outperform verbose alternatives
- Smaller models are more sensitive to prompt engineering
- Logic deviations are the dominant failure mode
- Reasoning strategy and output format significantly affect accuracy

Future Directions:

- Extend evaluation to more complex problems (difficulty 3-10)
- Test multilingual code generation across Python, Java, C++
- Investigate few-shot prompting with example solutions
- Develop automated prompt optimization frameworks